

Deep Learning

Ian Goodfellow
Yoshua Bengio
Aaron Courville

Contents

Website	vii
Acknowledgments	viii
Notation	xi
1 Introduction	1
1.1 Who Should Read This Book?	8
1.2 Historical Trends in Deep Learning	11
I Applied Math and Machine Learning Basics	29
2 Linear Algebra	31
2.1 Scalars, Vectors, Matrices and Tensors	31
2.2 Multiplying Matrices and Vectors	34
2.3 Identity and Inverse Matrices	36
2.4 Linear Dependence and Span	37
2.5 Norms	39
2.6 Special Kinds of Matrices and Vectors	40
2.7 Eigendecomposition	42
2.8 Singular Value Decomposition	44
2.9 The Moore-Penrose Pseudoinverse	45
2.10 The Trace Operator	46
2.11 The Determinant	47
2.12 Example: Principal Components Analysis	48
3 Probability and Information Theory	53
3.1 Why Probability?	54

3.2	Random Variables	56
3.3	Probability Distributions	56
3.4	Marginal Probability	58
3.5	Conditional Probability	59
3.6	The Chain Rule of Conditional Probabilities	59
3.7	Independence and Conditional Independence	60
3.8	Expectation, Variance and Covariance	60
3.9	Common Probability Distributions	62
3.10	Useful Properties of Common Functions	67
3.11	Bayes' Rule	70
3.12	Technical Details of Continuous Variables	71
3.13	Information Theory	72
3.14	Structured Probabilistic Models	75
4	Numerical Computation	80
4.1	Overflow and Underflow	80
4.2	Poor Conditioning	82
4.3	Gradient-Based Optimization	82
4.4	Constrained Optimization	93
4.5	Example: Linear Least Squares	96
5	Machine Learning Basics	98
5.1	Learning Algorithms	99
5.2	Capacity, Overfitting and Underfitting	110
5.3	Hyperparameters and Validation Sets	120
5.4	Estimators, Bias and Variance	122
5.5	Maximum Likelihood Estimation	131
5.6	Bayesian Statistics	135
5.7	Supervised Learning Algorithms	139
5.8	Unsupervised Learning Algorithms	145
5.9	Stochastic Gradient Descent	150
5.10	Building a Machine Learning Algorithm	152
5.11	Challenges Motivating Deep Learning	154
II	Deep Networks: Modern Practices	165
6	Deep Feedforward Networks	167
6.1	Example: Learning XOR	170
6.2	Gradient-Based Learning	176

6.3	Hidden Units	190
6.4	Architecture Design	196
6.5	Back-Propagation and Other Differentiation Algorithms	203
6.6	Historical Notes	224
7	Regularization for Deep Learning	228
7.1	Parameter Norm Penalties	230
7.2	Norm Penalties as Constrained Optimization	237
7.3	Regularization and Under-Constrained Problems	239
7.4	Dataset Augmentation	240
7.5	Noise Robustness	242
7.6	Semi-Supervised Learning	244
7.7	Multi-Task Learning	245
7.8	Early Stopping	246
7.9	Parameter Tying and Parameter Sharing	251
7.10	Sparse Representations	253
7.11	Bagging and Other Ensemble Methods	255
7.12	Dropout	257
7.13	Adversarial Training	267
7.14	Tangent Distance, Tangent Prop, and Manifold Tangent Classifier	268
8	Optimization for Training Deep Models	274
8.1	How Learning Differs from Pure Optimization	275
8.2	Challenges in Neural Network Optimization	282
8.3	Basic Algorithms	294
8.4	Parameter Initialization Strategies	301
8.5	Algorithms with Adaptive Learning Rates	306
8.6	Approximate Second-Order Methods	310
8.7	Optimization Strategies and Meta-Algorithms	318
9	Convolutional Networks	331
9.1	The Convolution Operation	332
9.2	Motivation	336
9.3	Pooling	340
9.4	Convolution and Pooling as an Infinitely Strong Prior	346
9.5	Variants of the Basic Convolution Function	348
9.6	Structured Outputs	359
9.7	Data Types	361
9.8	Efficient Convolution Algorithms	363
9.9	Random or Unsupervised Features	364

9.10	The Neuroscientific Basis for Convolutional Networks	365
9.11	Convolutional Networks and the History of Deep Learning	372
10	Sequence Modeling: Recurrent and Recursive Nets	374
10.1	Unfolding Computational Graphs	376
10.2	Recurrent Neural Networks	379
10.3	Bidirectional RNNs	396
10.4	Encoder-Decoder Sequence-to-Sequence Architectures	397
10.5	Deep Recurrent Networks	399
10.6	Recursive Neural Networks	401
10.7	The Challenge of Long-Term Dependencies	403
10.8	Echo State Networks	406
10.9	Leaky Units and Other Strategies for Multiple Time Scales	409
10.10	The Long Short-Term Memory and Other Gated RNNs	411
10.11	Optimization for Long-Term Dependencies	415
10.12	Explicit Memory	419
11	Practical methodology	424
11.1	Performance Metrics	425
11.2	Default Baseline Models	428
11.3	Determining Whether to Gather More Data	429
11.4	Selecting Hyperparameters	430
11.5	Debugging Strategies	439
11.6	Example: Multi-Digit Number Recognition	443
12	Applications	446
12.1	Large Scale Deep Learning	446
12.2	Computer Vision	455
12.3	Speech Recognition	461
12.4	Natural Language Processing	464
12.5	Other Applications	480
III	Deep Learning Research	489
13	Linear Factor Models	492
13.1	Probabilistic PCA and Factor Analysis	493
13.2	Independent Component Analysis (ICA)	494
13.3	Slow Feature Analysis	496
13.4	Sparse Coding	499

13.5	Manifold Interpretation of PCA	502
14	Autoencoders	505
14.1	Undercomplete Autoencoders	506
14.2	Regularized Autoencoders	507
14.3	Representational Power, Layer Size and Depth	511
14.4	Stochastic Encoders and Decoders	512
14.5	Denoising Autoencoders	513
14.6	Learning Manifolds with Autoencoders	518
14.7	Contractive Autoencoders	524
14.8	Predictive Sparse Decomposition	526
14.9	Applications of Autoencoders	527
15	Representation Learning	529
15.1	Greedy Layer-Wise Unsupervised Pretraining	531
15.2	Transfer Learning and Domain Adaptation	539
15.3	Semi-Supervised Disentangling of Causal Factors	544
15.4	Distributed Representation	549
15.5	Exponential Gains from Depth	556
15.6	Providing Clues to Discover Underlying Causes	557
16	Structured Probabilistic Models for Deep Learning	561
16.1	The Challenge of Unstructured Modeling	562
16.2	Using Graphs to Describe Model Structure	566
16.3	Sampling from Graphical Models	583
16.4	Advantages of Structured Modeling	584
16.5	Learning about Dependencies	585
16.6	Inference and Approximate Inference	586
16.7	The Deep Learning Approach to Structured Probabilistic Models	587
17	Monte Carlo Methods	593
17.1	Sampling and Monte Carlo Methods	593
17.2	Importance Sampling	595
17.3	Markov Chain Monte Carlo Methods	598
17.4	Gibbs Sampling	602
17.5	The Challenge of Mixing between Separated Modes	602
18	Confronting the Partition Function	608
18.1	The Log-Likelihood Gradient	609
18.2	Stochastic Maximum Likelihood and Contrastive Divergence	610

18.3	Pseudolikelihood	618
18.4	Score Matching and Ratio Matching	620
18.5	Denoising Score Matching	622
18.6	Noise-Contrastive Estimation	623
18.7	Estimating the Partition Function	626
19	Approximate inference	634
19.1	Inference as Optimization	636
19.2	Expectation Maximization	637
19.3	MAP Inference and Sparse Coding	638
19.4	Variational Inference and Learning	641
19.5	Learned Approximate Inference	653
20	Deep Generative Models	656
20.1	Boltzmann Machines	656
20.2	Restricted Boltzmann Machines	658
20.3	Deep Belief Networks	662
20.4	Deep Boltzmann Machines	665
20.5	Boltzmann Machines for Real-Valued Data	678
20.6	Convolutional Boltzmann Machines	685
20.7	Boltzmann Machines for Structured or Sequential Outputs	687
20.8	Other Boltzmann Machines	688
20.9	Back-Propagation through Random Operations	689
20.10	Directed Generative Nets	694
20.11	Drawing Samples from Autoencoders	712
20.12	Generative Stochastic Networks	716
20.13	Other Generation Schemes	717
20.14	Evaluating Generative Models	719
20.15	Conclusion	721
	Bibliography	723
	Index	780

Website

www.deeplearningbook.org

This book is accompanied by the above website. The website provides a variety of supplementary material, including exercises, lecture slides, corrections of mistakes, and other resources that should be useful to both readers and instructors.